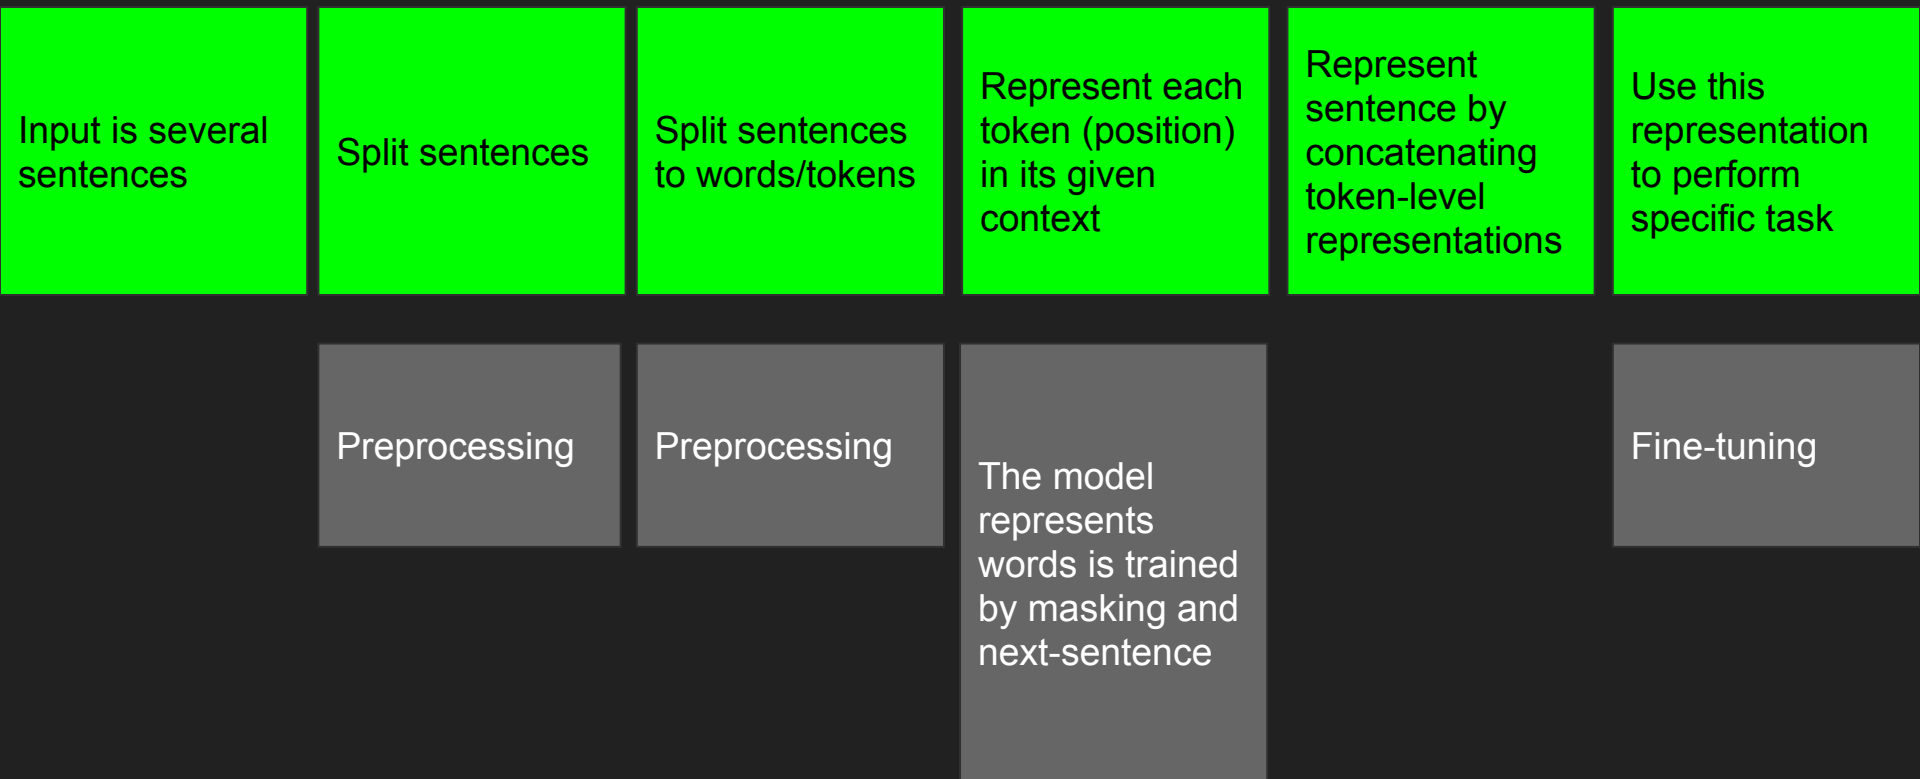


# Using BERT for Sentiment Analysis

# Outline

1. Justification
2. Preprocessing
3. Pretraining
4. Feature selection
5. Fine tuning
6. Understanding/Architecture
7. Regularization
8. What did not work(NBSVM)
9. Accuracies



# Outline

1. Justification
2. Preprocessing
3. Pretraining
4. Feature selection
5. Fine tuning
6. Understanding/Architecture
7. Regularization
8. What did not work(NBSVM)
9. Accuracies

# Data Preprocessing: Get data into format BERT understands

- Lowercase everything
- Tokenizes sentence to words, then words to WordPieces
  - (“calling” => [“call”, “##ing”])
- WordPieces are within corpus of 30,000 tokens
  - Constrained vocabulary of units so it's easier to process on
  - Handles rare words more effectively
- Add “CLS” and “SEP” tokens

# Pretraining

- Data: BooksCorpus and English Wikipedia (massive continuous text)
- Fine contextual understanding (word-level)
- Coarse contextual understanding (sentence-level)

# Pretraining, Fine

- Consider “bank” in “river bank” and “bank account”
  - Word2Vec represents them the same way (context-free)
  - BERT uses contextual representations of words
- Train model to predict words based on *entire general* context (bidirectional, random masking) instead of the *previous specific* context (left-directional)
- 15% chance a word gets chosen

- 80% of the time: Replace the word with the [MASK] token, e.g., my dog is hairy → my dog is [MASK]
- 10% of the time: Replace the word with a random word, e.g., my dog is hairy → my dog is apple
- 10% of the time: Keep the word unchanged, e.g., my dog is hairy → my dog is hairy. The purpose of this is to bias the representation towards the actual observed word.

# Pretraining, Coarse

- Train model to predict if some sentence is *next* sentence given *current* sentence

Input = [CLS] the man went to [MASK] store [SEP]

he bought a gallon [MASK] milk [SEP]

Label = IsNext

Input = [CLS] the man [MASK] to the store [SEP]

penguin [MASK] are flight ##less birds [SEP]

Label = NotNext



# Transformers

- Recurrence allows system to encode what was read before to better understand what is being read now
  - However, has bad long-range dependency, is greatly biased towards remembering words it saw most recently
- Use attention to **choose** which of prior words are **most important** to understanding the context as we continue reading
  - Far less bias on words that were just read, better long-range dependency
- Key Idea: Multi-Headed Block
  - One attention block captures some global context via weighted sum of hidden states
  - Use multiple attention blocks to capture many diff. contexts

# Transformers High Level

- Machine Translation uses basic-encoder decoder model:
  - Map sequences to sequences via intermediate representation. Maximize probability of correct prediction at each step
- Intuitively the attention mechanism allows the decoder to “look back” at entire sequence sentence
- Attention gives encoder access to all hidden states
- Encoder focuses and ignores hidden states by using weights
  - Decoder passed weighted sum to predict next word
  - Computed originally w/ feed forward neural network
- 3 types of dependencies
  - Input and output
  - Input
  - output

# Transformers low level

- Novelty of transformer is to allow decoder to see entire input sequence instead of left to right
- Multi-Head
  - Encoder hidden states(valued)
  - Decoder hidden state(query)
  - Block computes multiple attention weighted sums over values
    - Uses different linear transformations
- Scaled Dot Product Attention:
  - $\text{Attention}(Q,K,V) = \text{softmax}(QK^T/\sqrt{d_k})V$
  - Transformer rescales dot product to prevent from exploding to huge values
-

# Training - Fine-Tuning

- Add a output layer to overall BERT architecture so that output is in correct format for task
- Train whole model for specific task on associated input/output using pre-trained BERT parameters as initial parameters
- Modify batch size, learning rate, # of training epochs for task
  - Hyperparameters found using a search method for specific task
- Our case we use:
  - Softmax to get probabilities of each outcome
  - Dropout with 90% keep rate to prevent overfitting


# Results

- Project 1 training data accuracy
  - Positive: 0.89 on 25,000 instances
  - Negative: 0.89 on 25,000 instances
  - Total: 0.89 on 25,000 instances
- Project 1 testing accuracy: 86.08%
- Project 1 hidden accuracy: 86.67%
- Project 2 accuracy: 74.39%

# Bidirectional Encoder Representations from Transformers

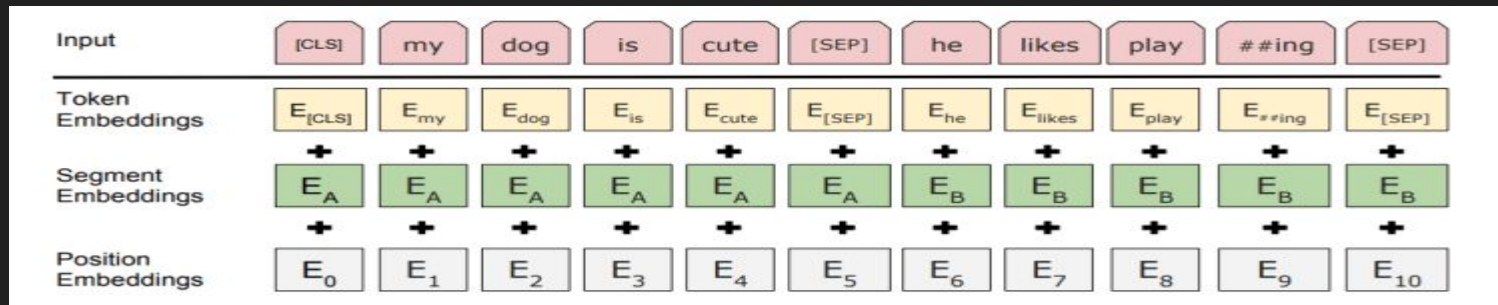
- Two main types of applying generic NLP models to specific tasks:
  - Feature based: Contextual representations of words as additional **features** (i.e. ELMo), task specific architectures
  - Fine-tuning based: Fine-tuning with few additional parameters, tuning of *all* parameters for task
- Both use unidirectional language models to learn representations
  - A word is represented by taking context of words that come before (but not after)

- BERT is a fine-tuning model
- Uses bi-directional representation of words
  - Word is represented by the context it appears in relative to words that come before **and** after
- This is in contrast to other models which generally only use information from right to left (instead of bidirectionally)



# Preprocessing

- As input BERT takes token embeddings as well as additional meta-embeddings
- BERT addresses transformed positional lacking by using positional embeddings
- Takes segment embeddings as input(pairs of sentence to improve contextualization)





# Masked Language Model

- Naively replace percentage with mask token and try to predict masked token
- BERT usually does this with 15%
- Additionally some random swapping of words done (~10%) of selected tokens to try to predict correctness regardless of what token is present

# Next Sentence Prediction Training

- Done for sentence relation training. Sentences separated by [SEP] token.  
50% of time second sentence is correct successor
- BERT must predict whether second sentence random

# Fine Tuning

- BERT encoder produces sequence of hidden states
  - Need to be reduced to single vector for classification
- BERT just takes hidden state corresponding to first token
- Sentence representation can then be fed into any classifier, can be tuned w/ r on top of BERT

# Hyperparameters

- Dropout: .1
- Batch: 32, 16
- Optimizer: Adam
- LR: 5e-5, 3e-5, 2e-5
- Epochs: 3, 4

## Pretraining:

- Seq len: 256
- Batch: 512

# Transformer Architecture

- RNNs used for a long time to do machine translation(mapping sentences to sentences). Problematic because they do not do well with large texts.  
Bottlenecked in intermediate representation
- Transformer composed of multiple attention blocks
- Stacks a layer that maps sequences to sequences
- Cannot take order of inputs into account(will treat the first and last tokens of inputs same if same word)
- BERT addressed by